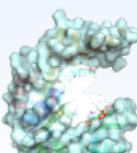


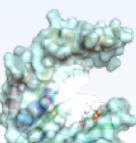


CASE STUDIES



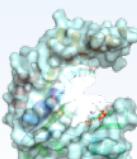
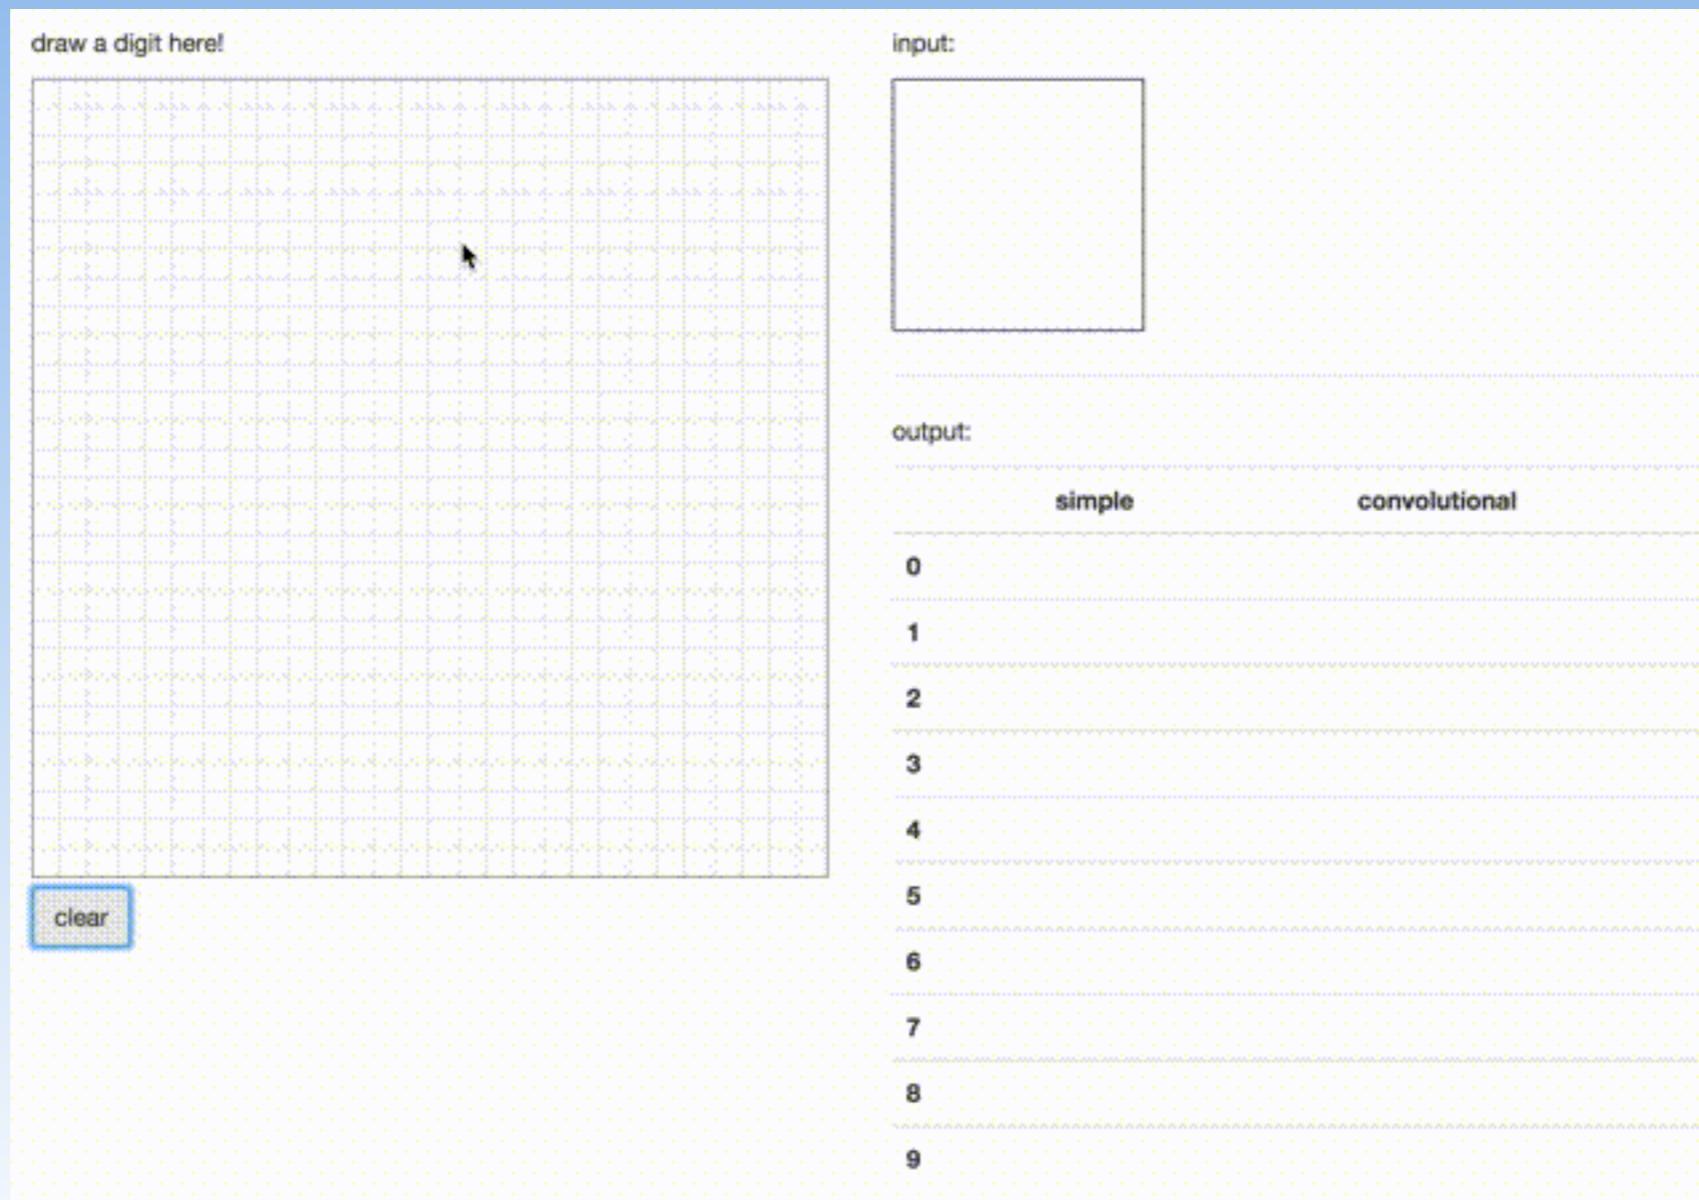
CASE STUDY

TENSORFLOW-MNIST



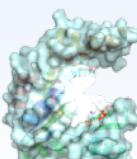
Catalit LLC

TENSORFLOW-MNIST



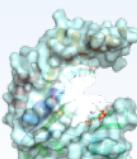
Catalit LLC

```
8  x = tf.placeholder("float", [None, 784])
9  sess = tf.Session()
10
11 # restore trained data
12 with tf.variable_scope("regression"):
13     y1, variables = model.regression(x)
14 saver = tf.train.Saver(variables)
15 saver.restore(sess, "mnist/data/regression.ckpt")
16
17
18 with tf.variable_scope("convolutional"):
19     keep_prob = tf.placeholder("float")
20     y2, variables = model.convolutional(x, keep_prob)
21 saver = tf.train.Saver(variables)
22 saver.restore(sess, "mnist/data/convolutional.ckpt")
23
24
25 def regression(input):
26     return sess.run(y1, feed_dict={x: input}).flatten().tolist()
27
28
29 def convolutional(input):
30     return sess.run(y2, feed_dict={x: input, keep_prob: 1.0}).flatten().tolist()
```



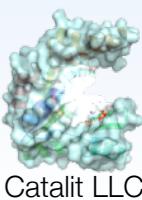
CODE

```
33 # webapp
34 app = Flask(__name__)
35
36
37 @app.route('/api/mnist', methods=['POST'])
38 def mnist():
39     input = ((255 - np.array(request.json, dtype=np.uint8)) / 255.0).reshape(1, 784)
40     output1 = regression(input)
41     output2 = convolutional(input)
42     return jsonify(results=[output1, output2])
43
44
45 @app.route('/')
46 def main():
47     return render_template('index.html')
48
49
50 if __name__ == '__main__':
51     app.run()
```



SUMMARY

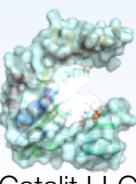
- WHY IS IT IMPORTANT?
 - model save / restore
 - flask serving of model prediction



RUN IT!

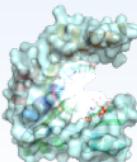
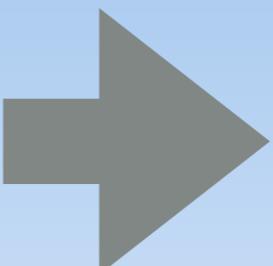
- source activate dwdlagpu
- cd 6_case_studies/tensorflow-mnist
- ./run.sh
- point browser to: <http://<your-ip>:8000>

CASE STUDY NEURAL STYLE TRANSFER



Catalit LLC

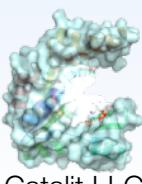
NEURAL STYLE TRANSFER



Catalit LLC

SUMMARY

- WHY IS IT IMPORTANT?
 - custom loss
 - uses pre-trained VGG model

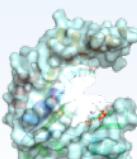


RUN IT!

- source activate dwdlagpu
- cd 6_case_studies/neural_style_transfer
- python neural_style_transfer.py img/obama.jpg img/starry_night.jpg results/obamagogh
- to download:
 - from laptop do:
 - scp -r -i dataweekends.pem ubuntu@<your-ip>:~/DL_Advanced_Workshop-03-11-2017/6_case_studies/neural_style_transfer/results .

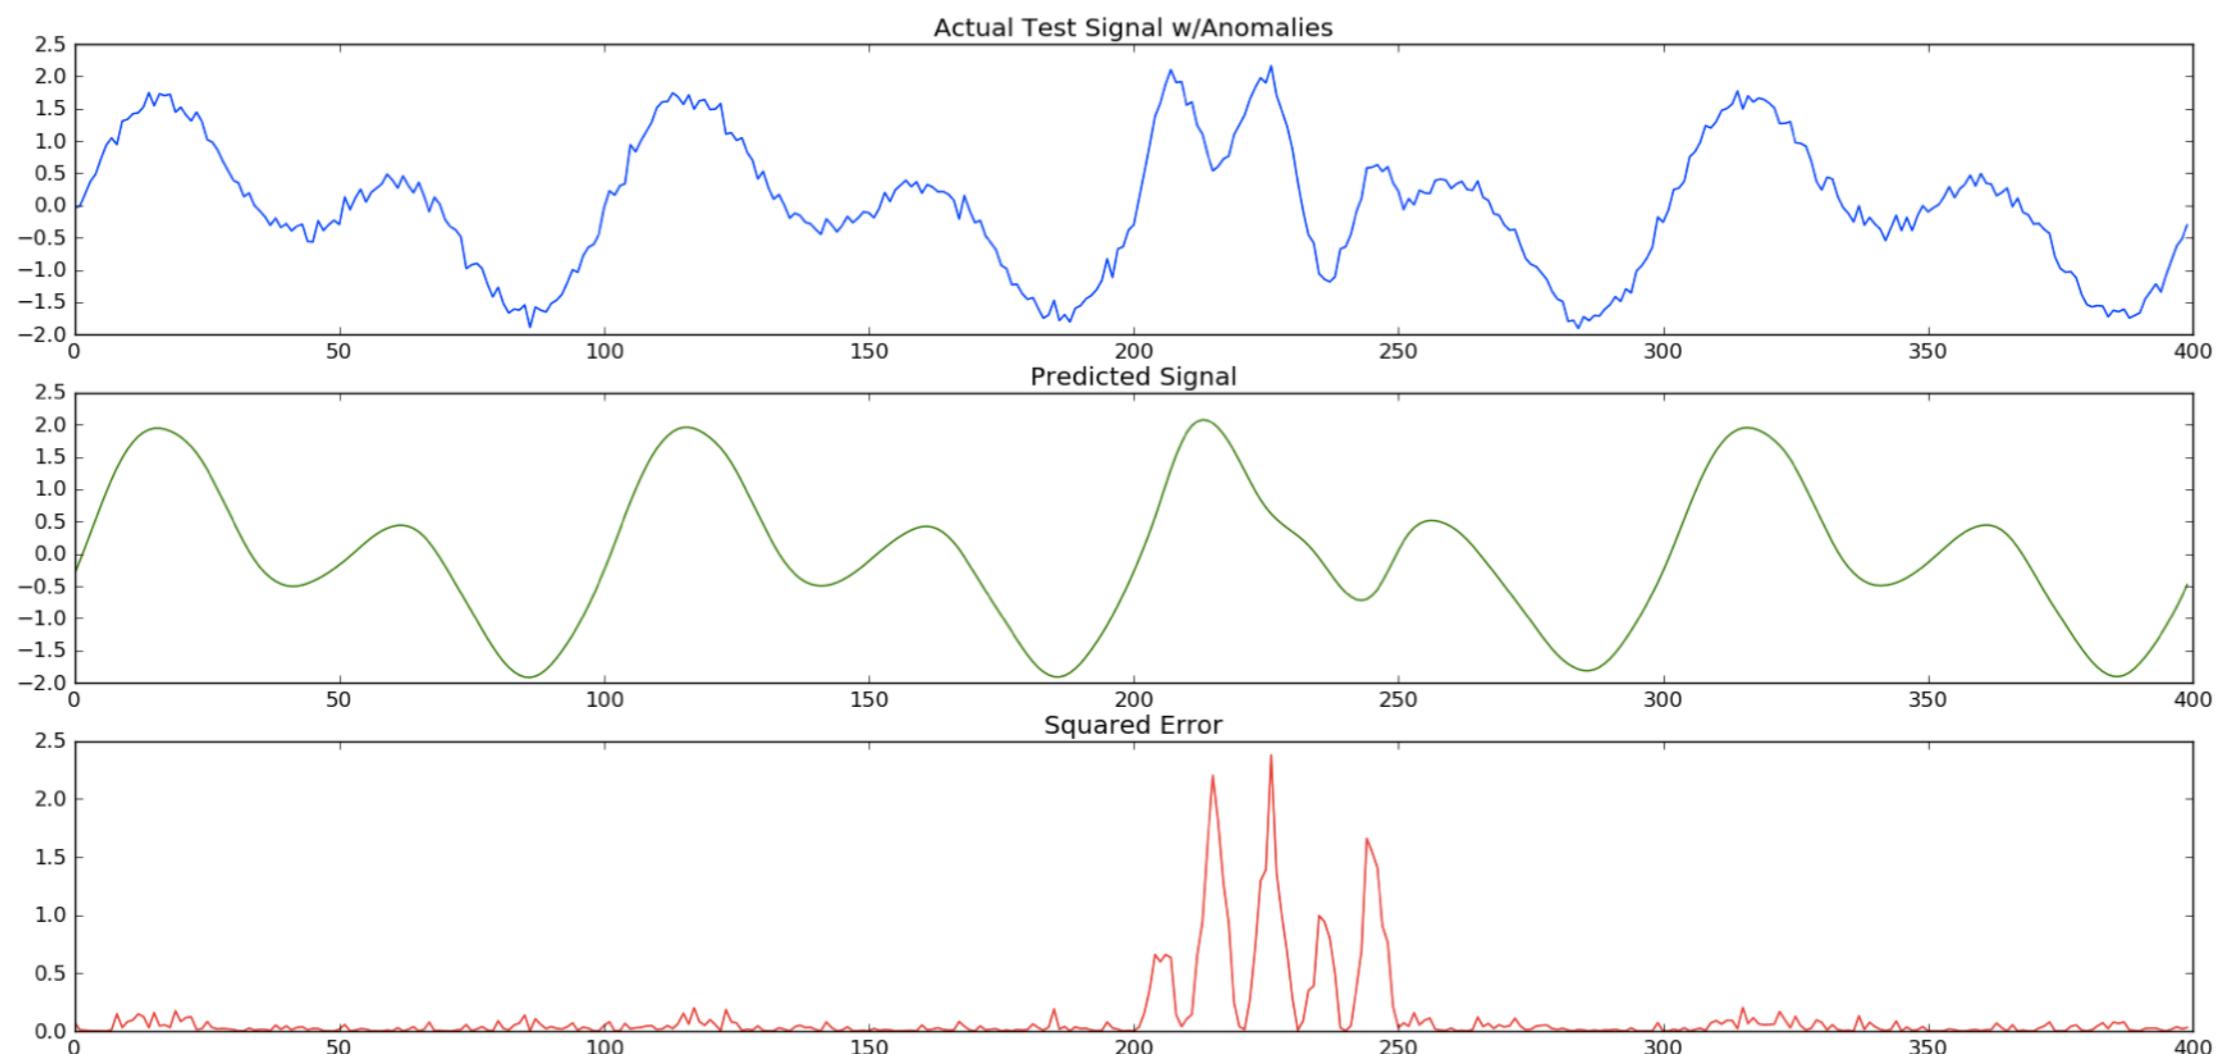
CASE STUDY

LSTM ANOMALY DETECTION

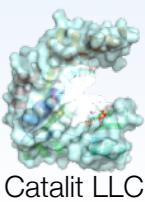


Catalit LLC

LSTM-ANOMALY

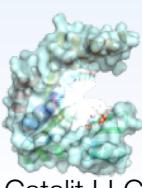


<https://github.com/ghego/lstm-anomaly-detect>



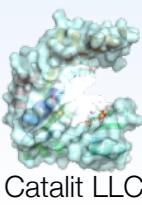
CODE

```
115 def build_model():
116     model = Sequential()
117     layers = {'input': 1, 'hidden1': 64, 'hidden2': 256, 'hidden3': 100, 'output': 1}
118
119     model.add(LSTM(
120         input_length=sequence_length - 1,
121         input_dim=layers['input'],
122         output_dim=layers['hidden1'],
123         return_sequences=True))
124     model.add(Dropout(0.2))
125
126     model.add(LSTM(
127         layers['hidden2'],
128         return_sequences=True))
129     model.add(Dropout(0.2))
130
131     model.add(LSTM(
132         layers['hidden3'],
133         return_sequences=False))
134     model.add(Dropout(0.2))
135
136     model.add(Dense(
137         output_dim=layers['output']))
138     model.add(Activation("linear"))
139
140     start = time.time()
141     model.compile(loss="mse", optimizer="rmsprop")
142     print "Compilation Time : ", time.time() - start
143     return model
```



SUMMARY

- WHY IS IT IMPORTANT?
 - LSTM on time series
 - stacked LSTM

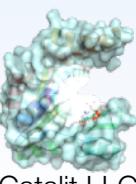


RUN IT!

- cd 6_case_studies/lstm-anomaly-detection
- python lstm-synthetic-wave-anomaly-detect.py

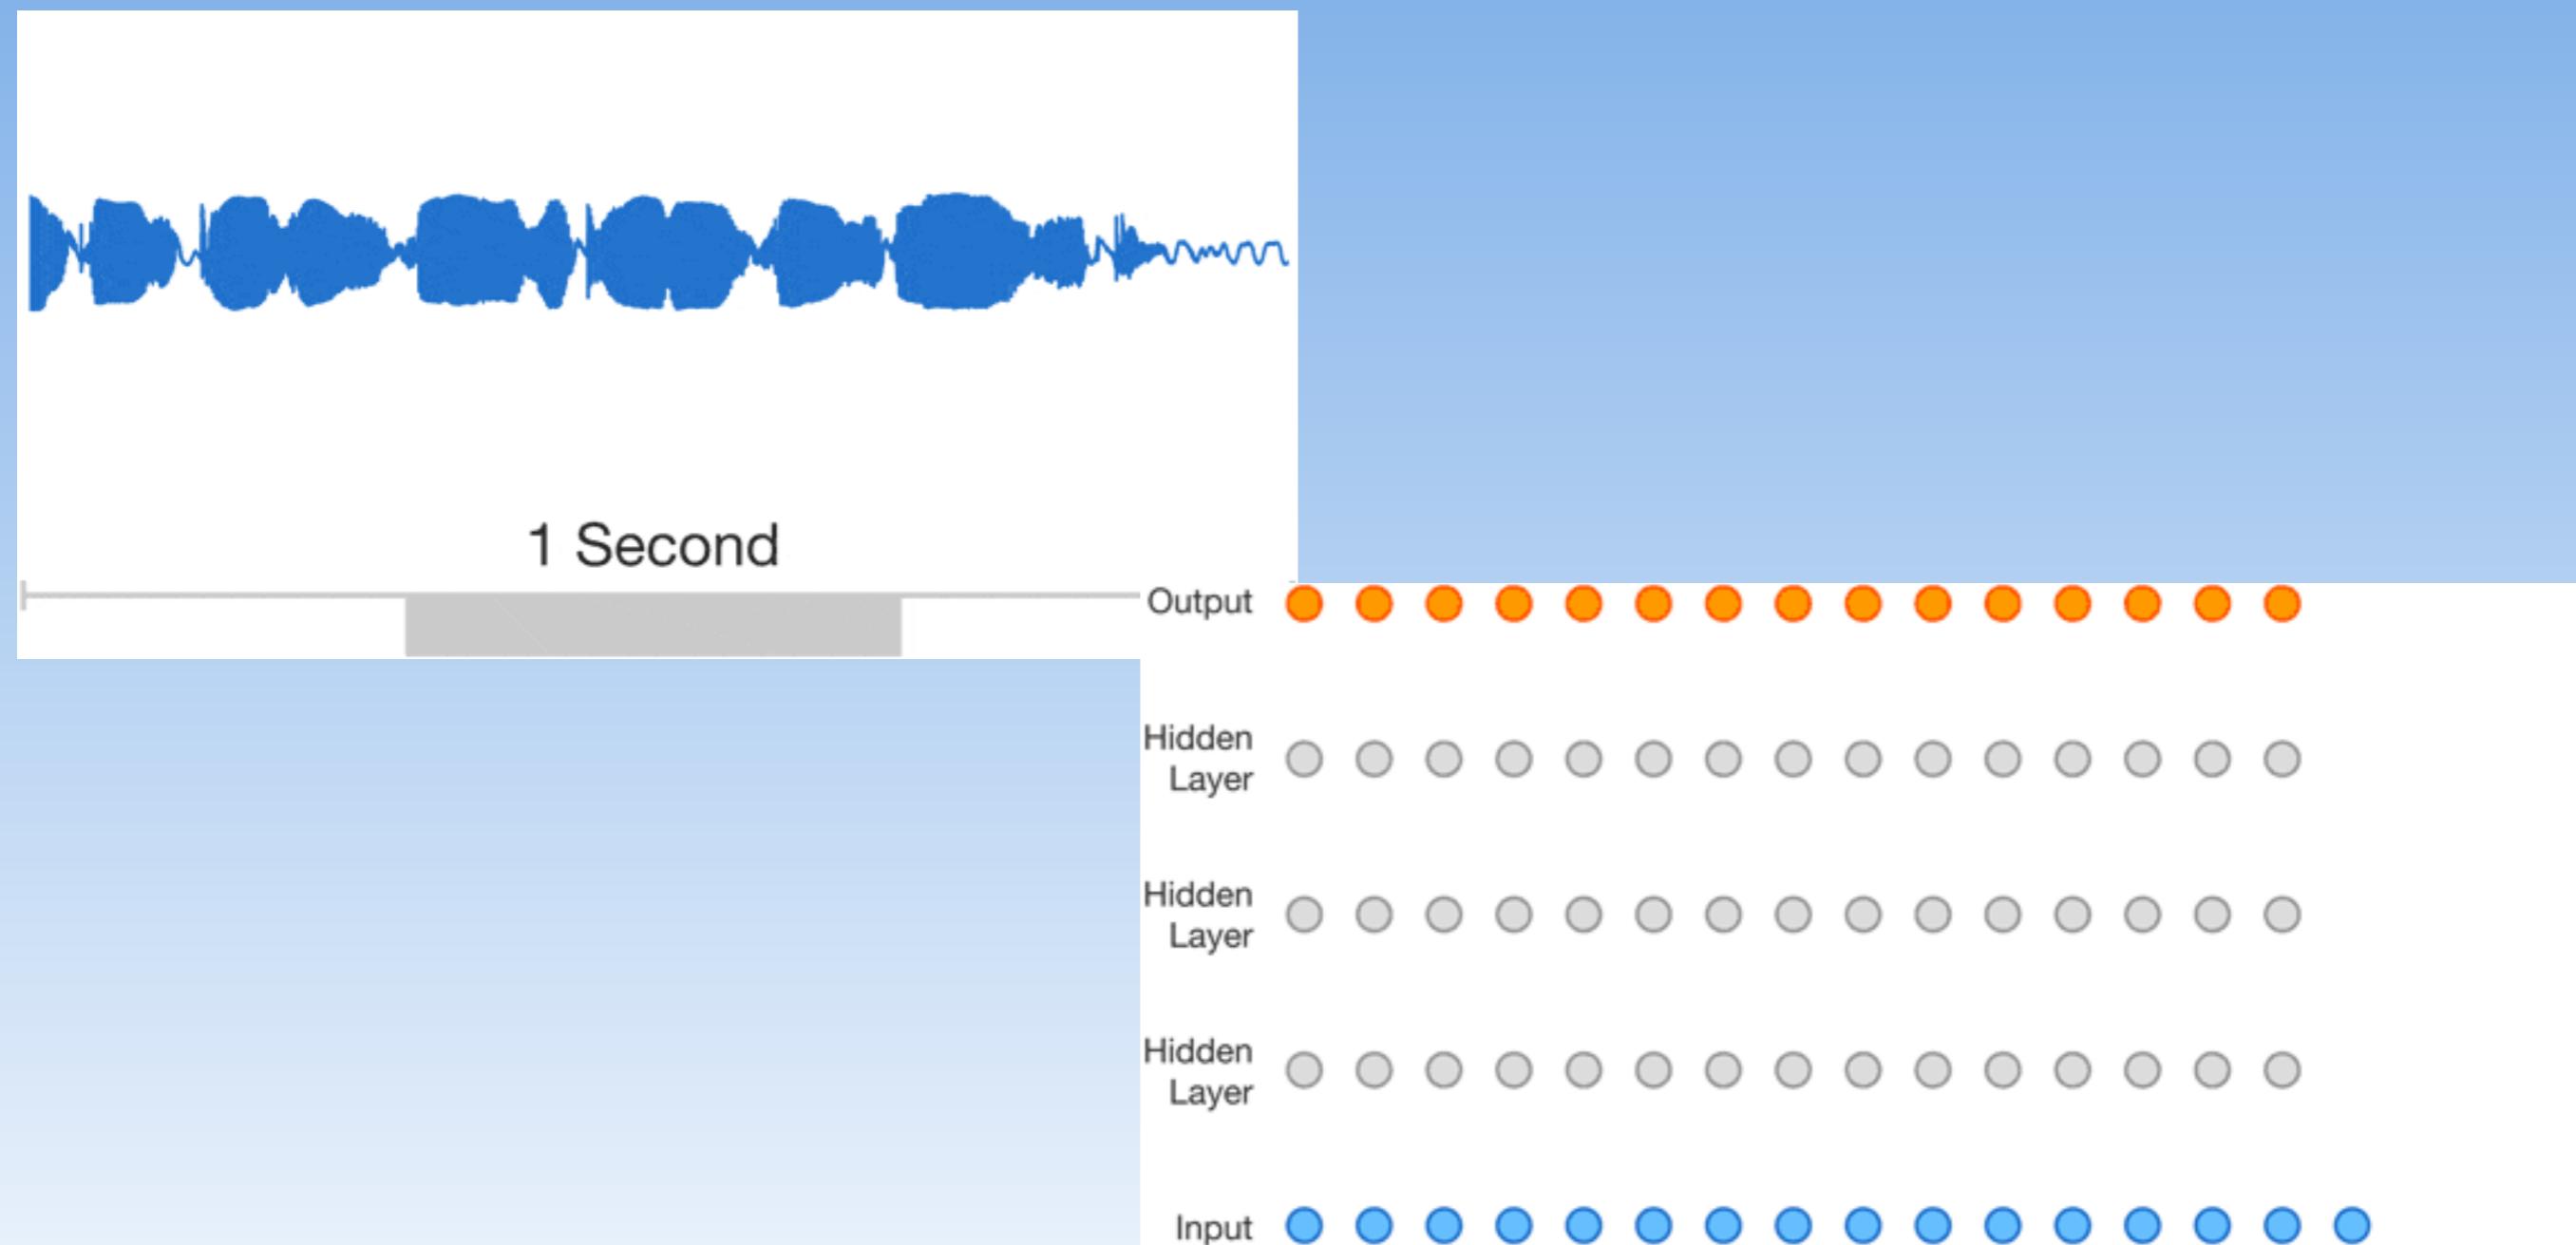
CASE STUDY

WAVENET



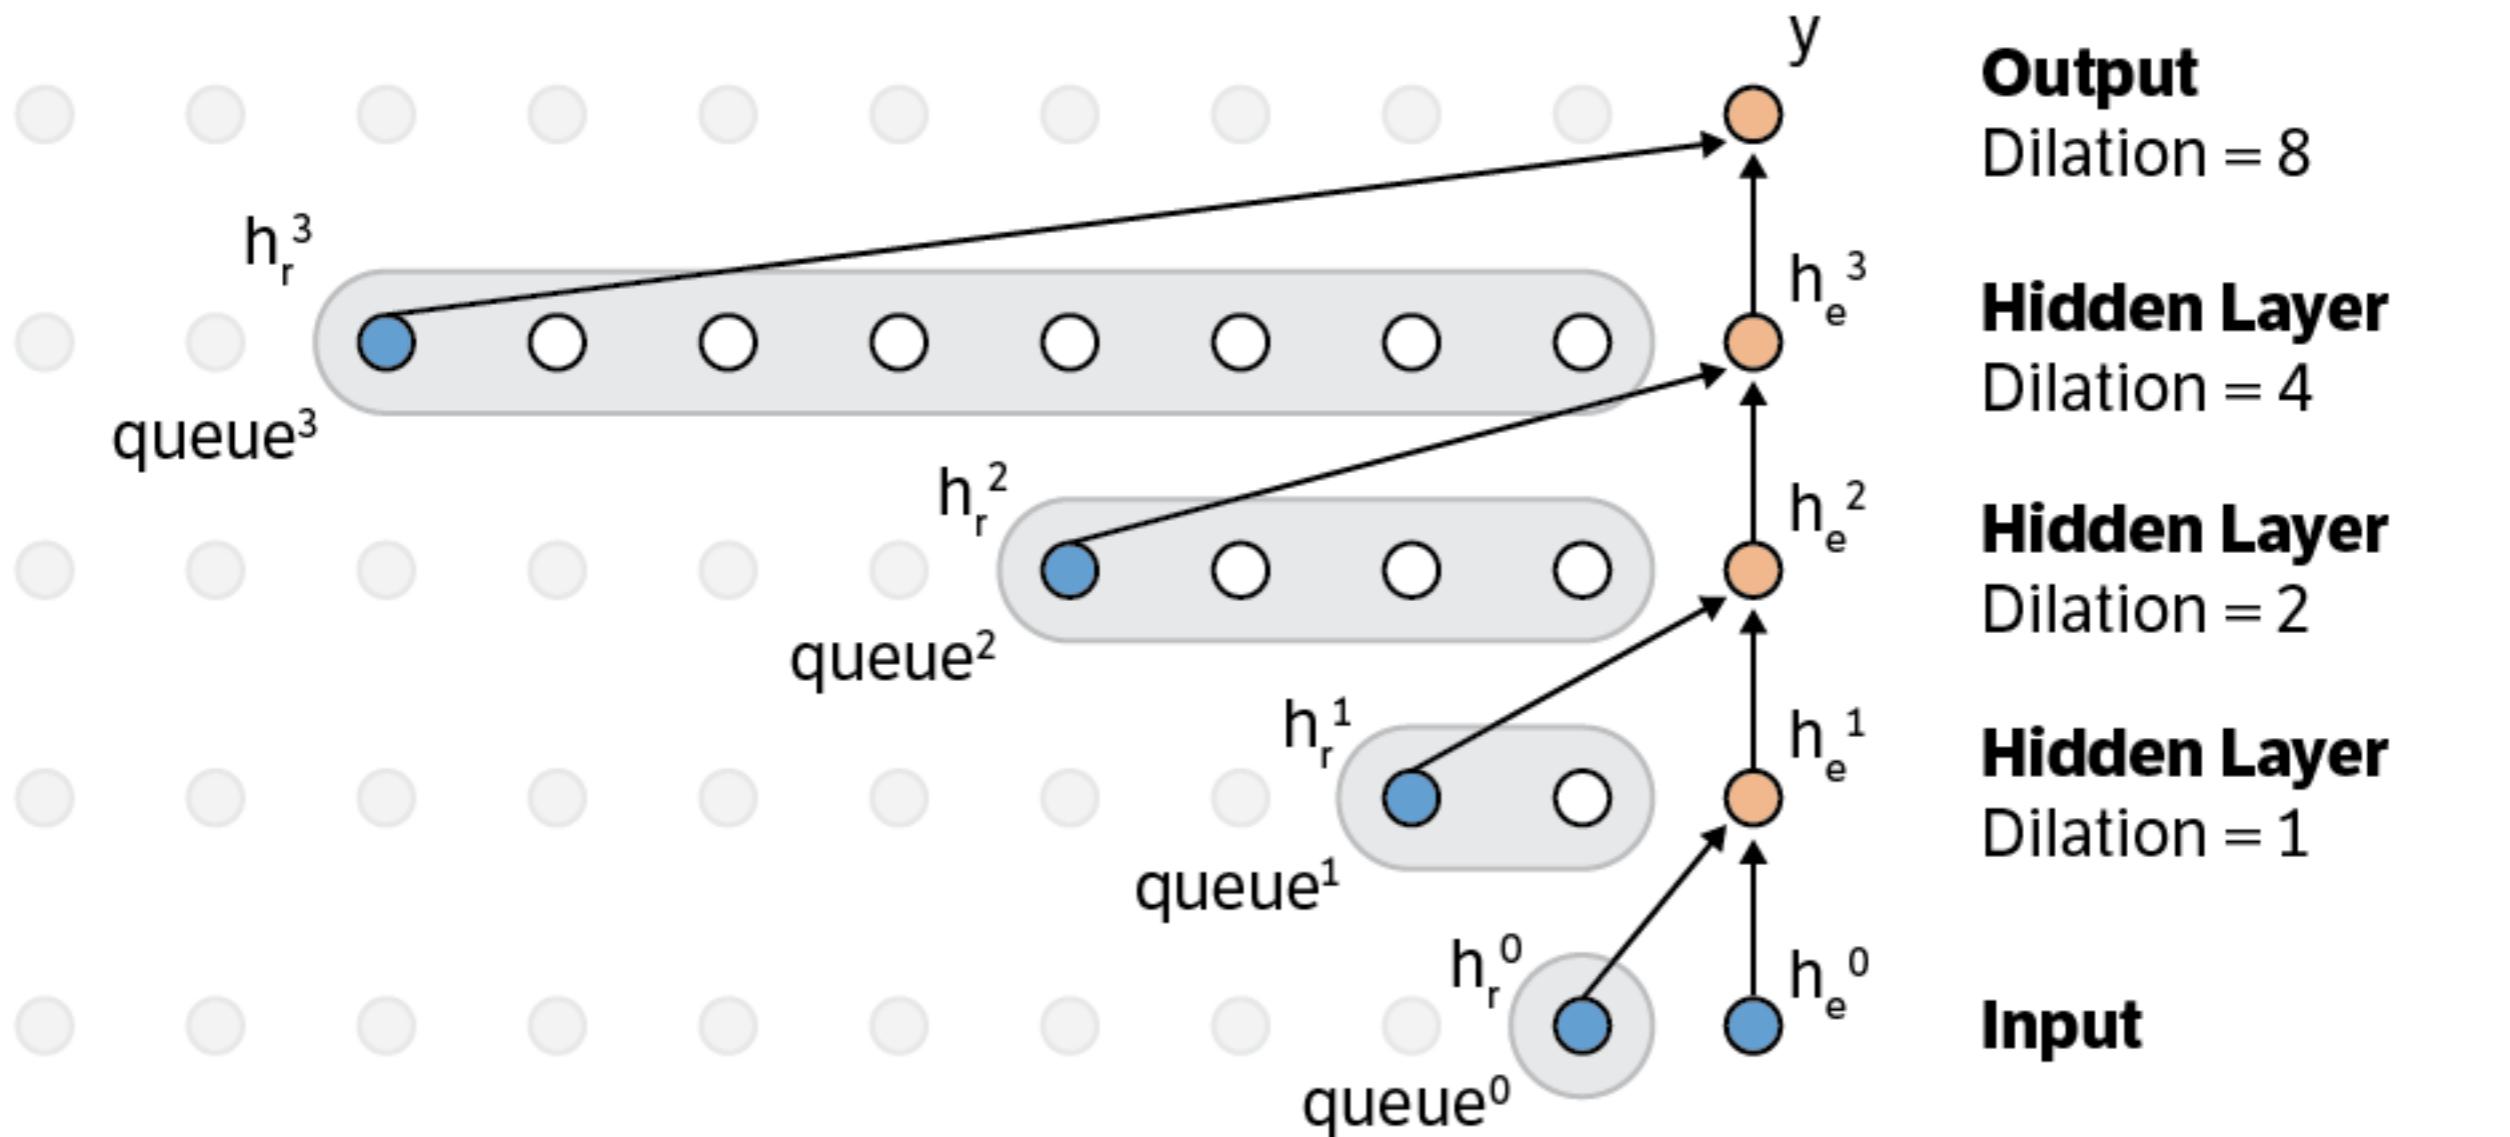
Catalit LLC

WAVENET



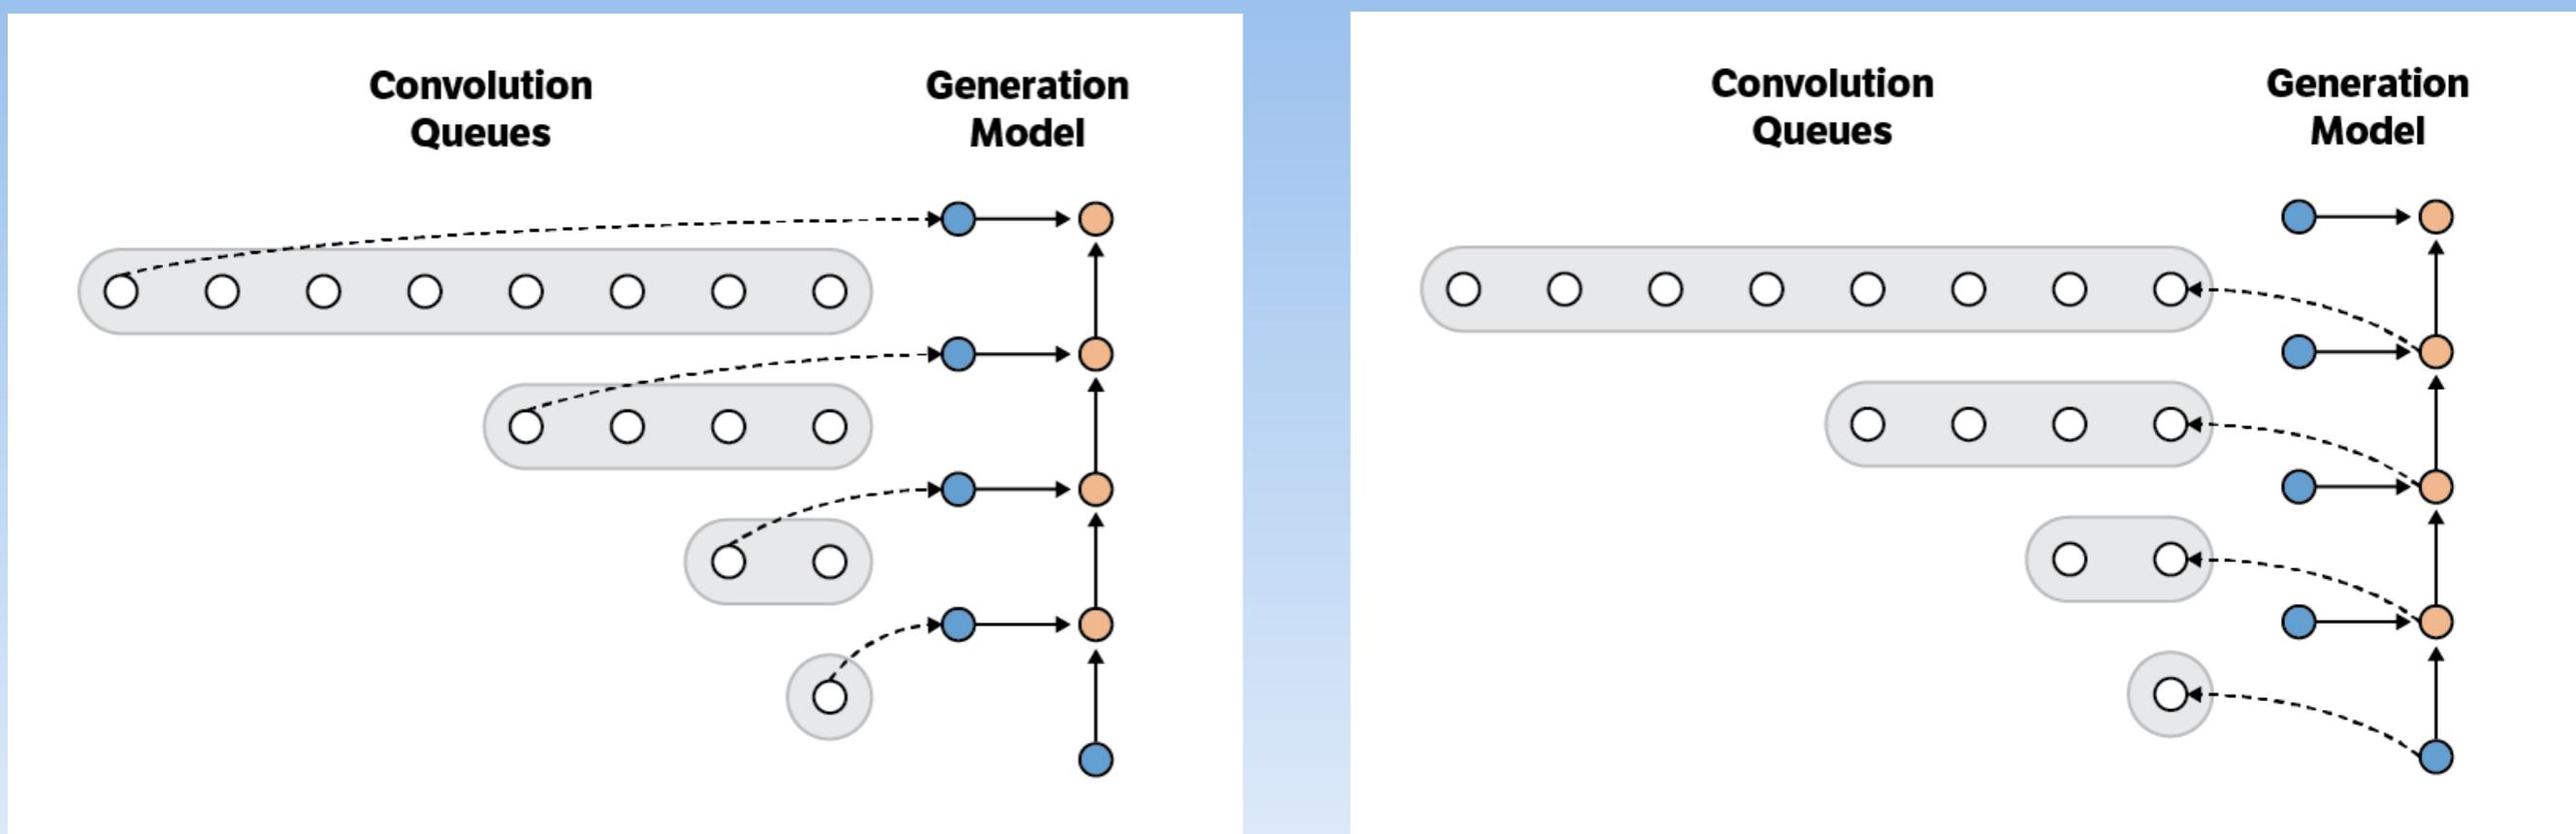
<https://deepmind.com/blog/wavenet-generative-model-raw-audio/>

FAST-WAVENET



<https://github.com/tomlepaine/fast-wavenet>

IMPLEMENTATION

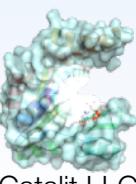


SUMMARY

- WHY IS IT IMPORTANT?
 - stacked convolution
 - convolutions with gaps
 - pixel level model

RUN IT!

- open jupyter notebook
- browse to folder 6_case_studies/fast-wavenet

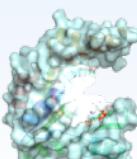


Catalit LLC

SEE ALSO

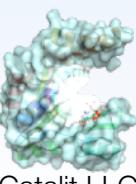
```
def wavenetBlock(n_atrous_filters, atrous_filter_size, atrous_rate):
    def f(input_):
        residual = input_
        tanh_out = AtrousConvolution1D(n_atrous_filters, atrous_filter_size,
                                         atrous_rate=atrous_rate,
                                         border_mode='same',
                                         activation='tanh')(input_)
        sigmoid_out = AtrousConvolution1D(n_atrous_filters, atrous_filter_size,
                                           atrous_rate=atrous_rate,
                                           border_mode='same',
                                           activation='sigmoid')(input_)
        merged = merge([tanh_out, sigmoid_out], mode='mul')
        skip_out = Convolution1D(1, 1, activation='relu', border_mode='same')(merged)
        out = merge([skip_out, residual], mode='sum')
        return out, skip_out
    return f
```

<https://github.com/usernaamee/keras-wavenet>



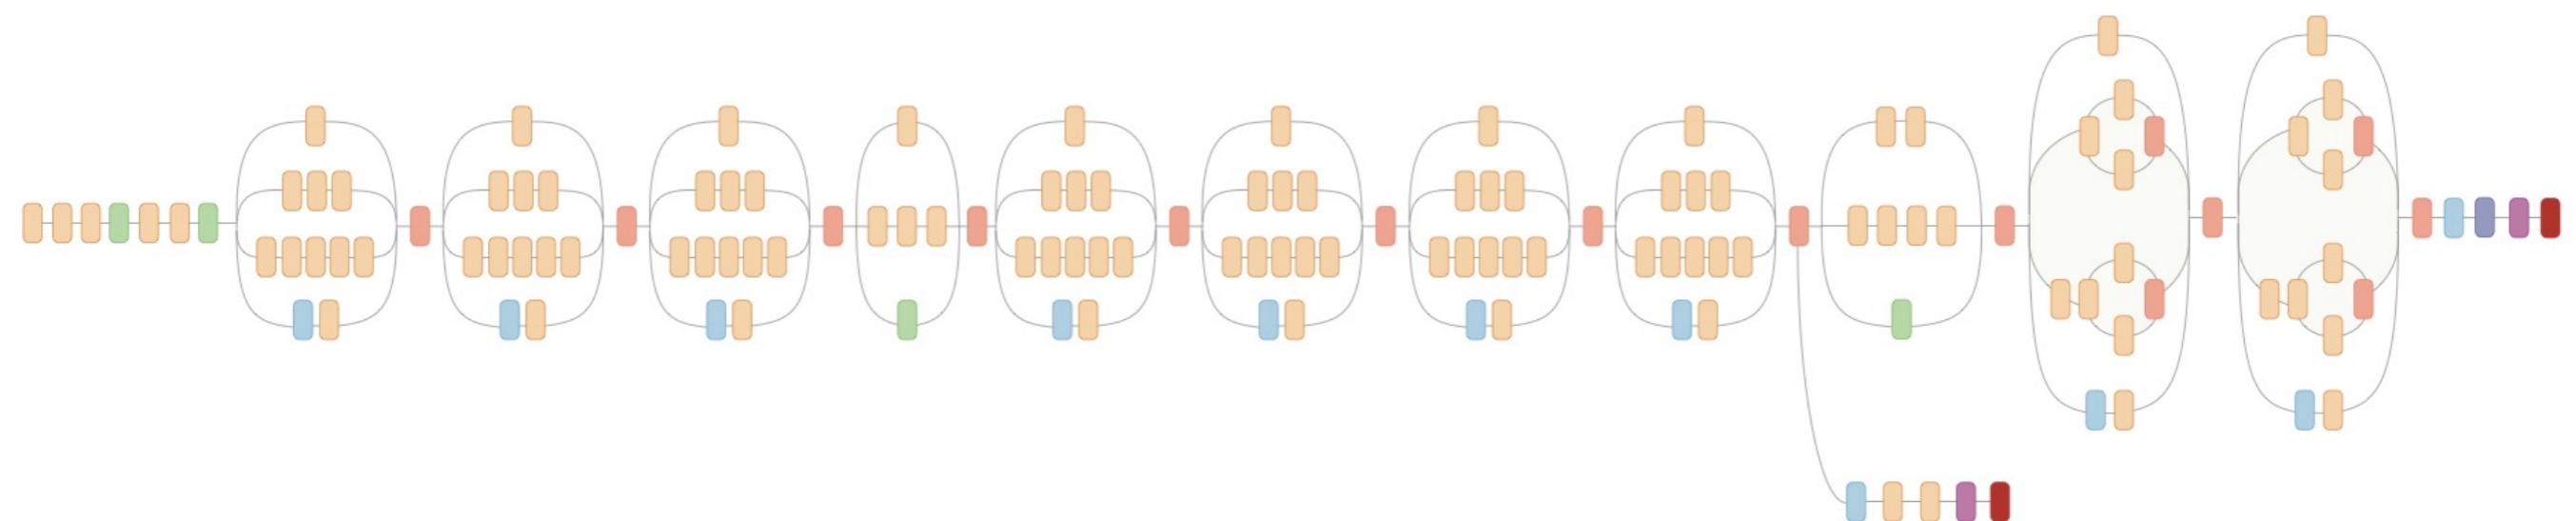
Catalit LLC

CASE STUDY INCEPTION RETRAIN



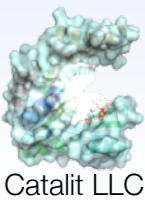
Catalit LLC

INCEPTION

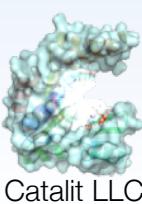


- Convolution
- AvgPool
- MaxPool
- Concat
- Dropout
- Fully connected
- Softmax

<http://arxiv.org/abs/1512.00567>



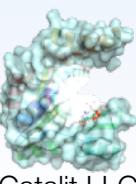
INCEPTION



Catalit LLC

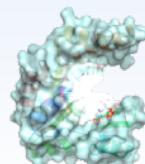
INCEPTION

RETRAIN :)



Catalit LLC

DETECT FLOWERS



Catalit LLC

RUN IT!

- cd 6_case_studies/inception-retrain
- ./retrain.sh
- python ./label_image.py ./tf_files/flower_photos/
daisy/21652746_cc379e0eea_m.jpg

CODE

```
DATA_URL = 'http://download.tensorflow.org/models/image/imagenet/inception-2015-12-05.tgz'  
# pylint: enable=line-too-long  
BOTTLENECK_TENSOR_NAME = 'pool_3/_reshape:0'  
BOTTLENECK_TENSOR_SIZE = 2048  
  
JPEG_DATA_TENSOR_NAME = 'DecodeJpeg/contents:0'  
RESIZED_INPUT_TENSOR_NAME = 'ResizeBilinear:0'  
MAX_NUM_IMAGES_PER_CLASS = 2 ** 27 - 1 # ~134M
```

```
def create_inception_graph():
```

"""Creates a graph from saved GraphDef file and returns a Graph object.

Returns:

Graph holding the trained Inception network, and various tensors we'll be manipulating.

"""

```
with tf.Session() as sess:
```

```
    model_filename = os.path.join(
```

```
        FLAGS.model_dir, 'classify_image_graph_def.pb')
```

```
    with gfile.FastGFile(model_filename, 'rb') as f:
```

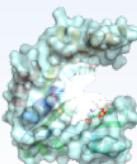
```
        graph_def = tf.GraphDef()
```

```
        graph_def.ParseFromString(f.read())
```

```
        bottleneck_tensor, jpeg_data_tensor, resized_input_tensor = (
```

```
            tf.import_graph_def(graph_def, name='', return_elements=[  
                BOTTLENECK_TENSOR_NAME, JPEG_DATA_TENSOR_NAME,  
                RESIZED_INPUT_TENSOR_NAME]))
```

```
    return sess.graph, bottleneck_tensor, jpeg_data_tensor, resized_input_tensor
```



Catalit LLC

BOTTLENECKS

```
def run_bottleneck_on_image(sess, image_data, image_data_tensor,  
                           bottleneck_tensor):  
    """Runs inference on an image to extract the 'bottleneck' summary layer.
```

Args:

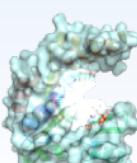
```
sess: Current active TensorFlow Session.  
image_data: String of raw JPEG data.  
image_data_tensor: Input data layer in the graph.  
bottleneck_tensor: Layer before the final softmax.
```

Returns:

```
Numpy array of bottleneck values.
```

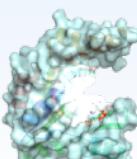
"""

```
bottleneck_values = sess.run(  
    bottleneck_tensor,  
    {image_data_tensor: image_data})  
bottleneck_values = np.squeeze(bottleneck_values)  
return bottleneck_values
```



RETRAINING

```
# Organizing the following ops as `final_training_ops` so they're easier
# to see in TensorBoard
layer_name = 'final_training_ops'
with tf.name_scope(layer_name):
    with tf.name_scope('weights'):
        layer_weights = tf.Variable(tf.truncated_normal(
            [BOTTLENECK_TENSOR_SIZE, class_count], stddev=0.001), name='final_weights')
        variable_summaries(layer_weights)
    with tf.name_scope('biases'):
        layer_biases = tf.Variable(
            tf.zeros([class_count]), name='final_biases')
        variable_summaries(layer_biases)
    with tf.name_scope('Wx_plus_b'):
        logits = tf.matmul(bottleneck_input, layer_weights) + layer_biases
        tf.summary.histogram('pre_activations', logits)
```

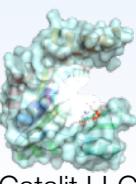


INCEPTION SUMMARY

- WHY IS IT IMPORTANT?
 - shows how to re-use a pre-trained network
 - piggyback on shoulders of giants
 - tensorflow name scopes
 - data augmentation
 - argument parsing

CASE STUDY

DEEP ANPR



Catalit LLC

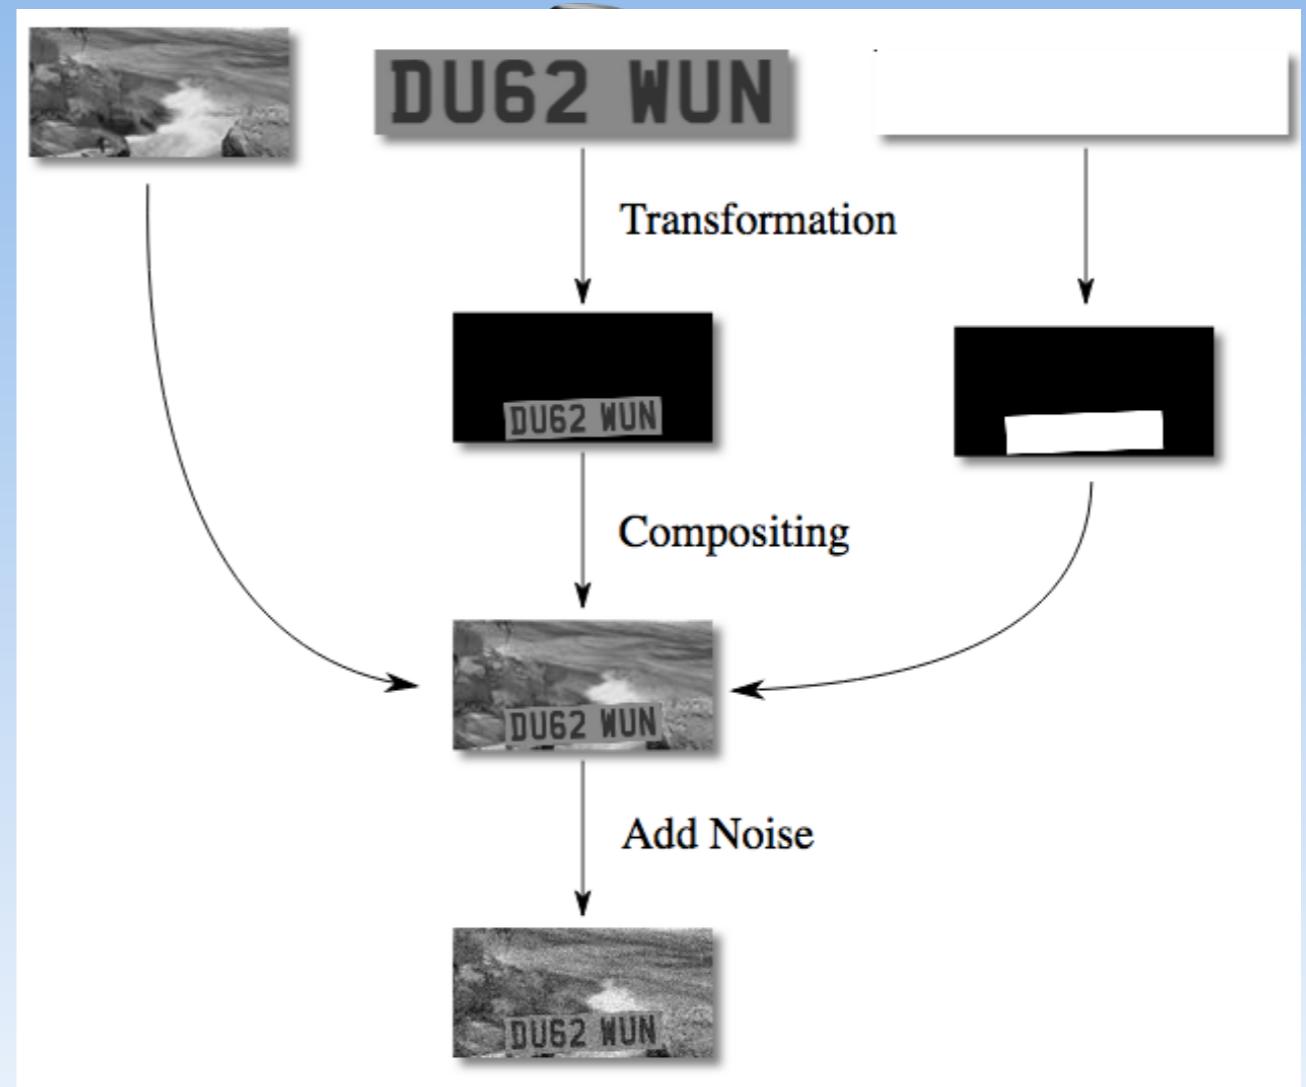
PLATE RECOGNITION WITH CNN



<http://matthewearl.github.io/2016/05/06/cnn-anpr/>

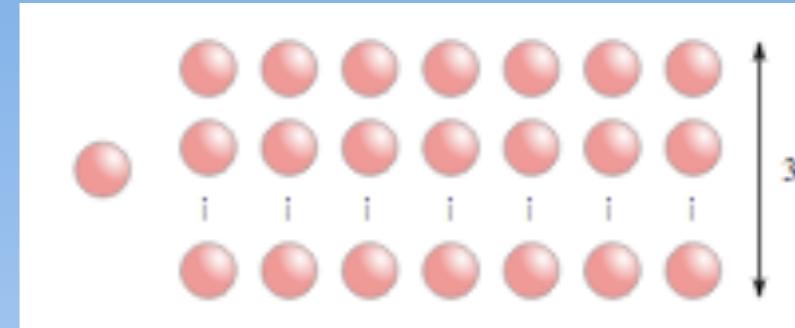
IMAGE SYNTHESIS

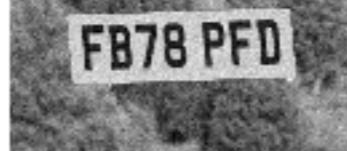
- synthetic images generated in the cpu during training
- using opencv

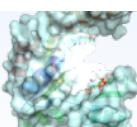


LOSS FUNCTION

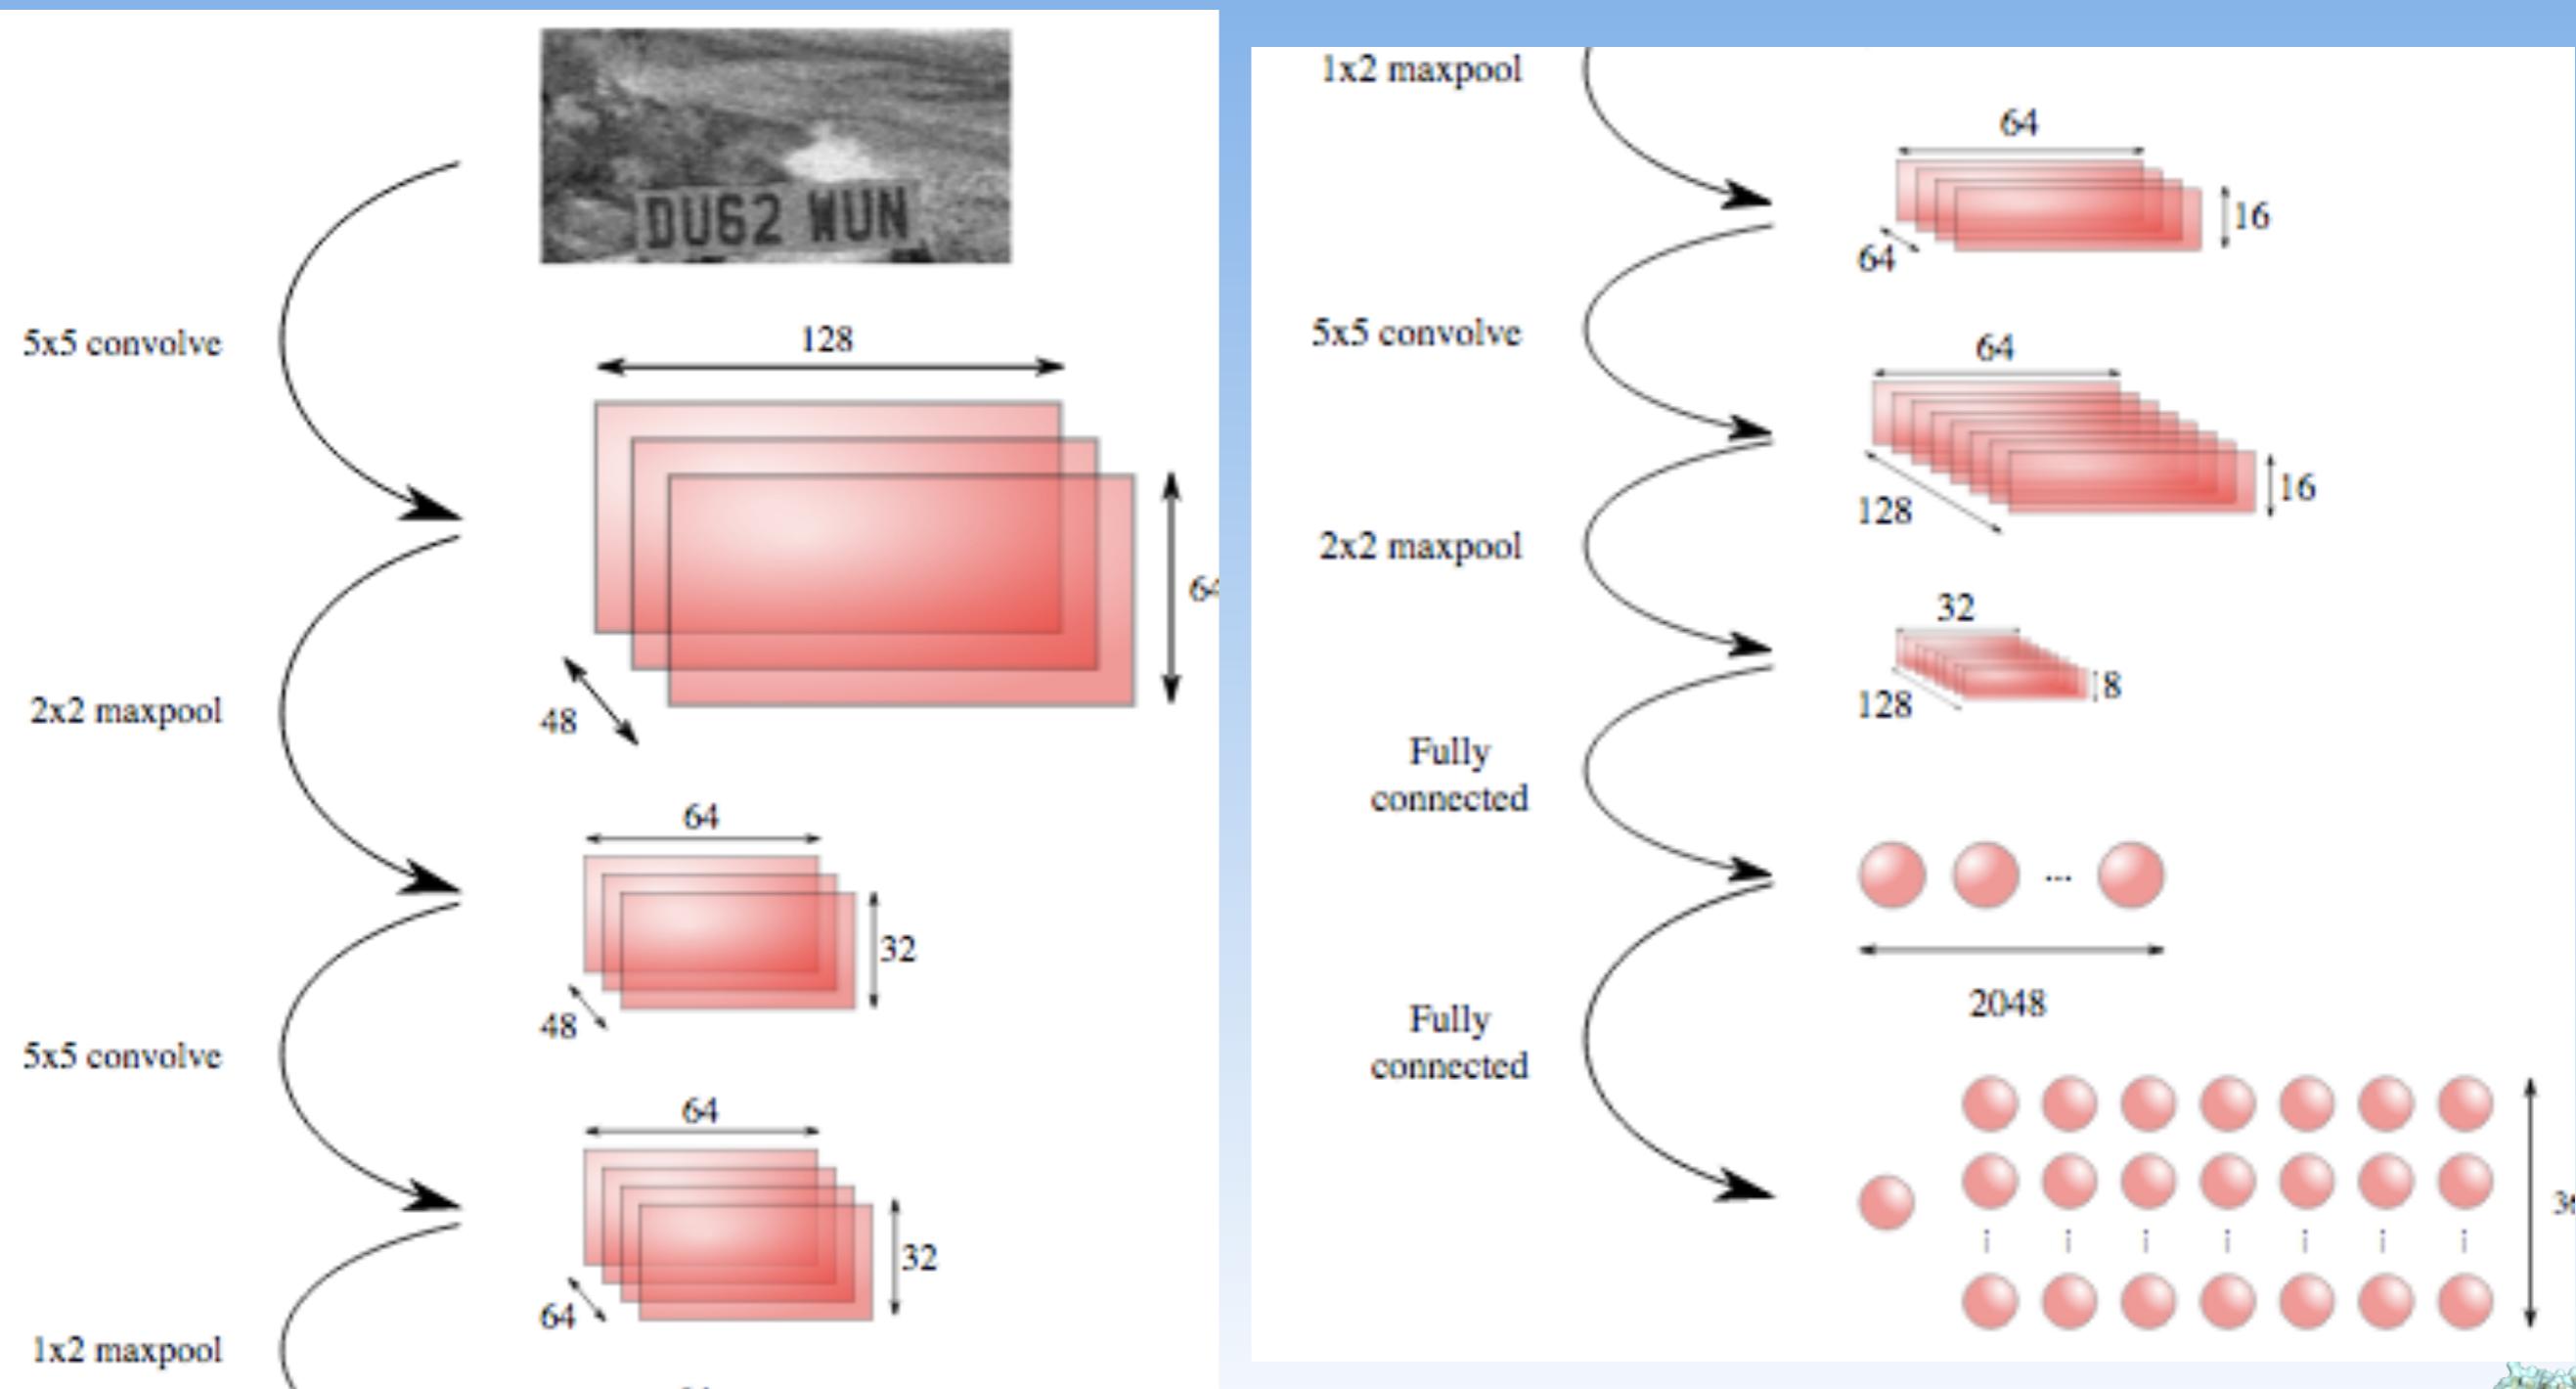
- 1 bit for presence
=> sigmoid
- 7 chars plate (out of 36)
=> softmax on each



-  expected output **HH41RFP 1**.
-  expected output **FB78PFD 1**.
-  expected output **JW01GAI 0**. (Plate partially truncated.)
-  expected output **AM46KVG 0**. (Plate too small.)
-  expected output **XG86KIO 0**. (Plate too big.)
-  expected output **XH07NYO 0**. (Plate not present at all.)



MODEL



PREDICTION

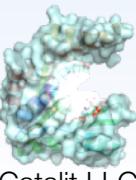
- get prediction at every window
- clever trick => map the fully connected part to convolutional layers



Catalit LLC

RUN IT!

- cd 6_case_studies/deep-anpr
- ./train.py
- when tired stop it
- ./detect.py test/<some_image>.png weights.npz
out.png



Catalit LLC

CODE

```
def generate_im(char_ims, num_bg_images):
    bg = generate_bg(num_bg_images)

    plate, plate_mask, code = generate_plate(FONT_HEIGHT, char_ims)

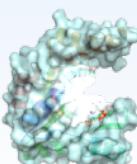
    M, out_of_bounds = make_affine_transform(
        from_shape=plate.shape,
        to_shape=bg.shape,
        min_scale=0.6,
        max_scale=0.875,
        rotation_variation=1.0,
        scale_variation=1.5,
        translation_variation=1.2)
    plate = cv2.warpAffine(plate, M, (bg.shape[1], bg.shape[0]))
    plate_mask = cv2.warpAffine(plate_mask, M, (bg.shape[1], bg.shape[0]))

    out = plate * plate_mask + bg * (1.0 - plate_mask)

    out = cv2.resize(out, (OUTPUT_SHAPE[1], OUTPUT_SHAPE[0]))

    out += numpy.random.normal(scale=0.05, size=out.shape)
    out = numpy.clip(out, 0., 1.)

    return out, code, not out_of_bounds
```



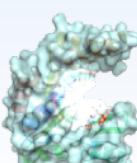
CODE

```
def get_loss(y, y_):
    # Calculate the loss from digits being incorrect.  Don't count loss from
    # digits that are in non-present plates.
    digits_loss = tf.nn.softmax_cross_entropy_with_logits(
        logits=tf.reshape(y[:, 1:],
                          [-1, len(common.CHARS)]),
        labels=tf.reshape(y_[:, 1:],
                          [-1, len(common.CHARS)]))

    digits_loss = tf.reshape(digits_loss, [-1, 7])
    digits_loss = tf.reduce_sum(digits_loss, 1)
    digits_loss *= (y_[:, 0] != 0)
    digits_loss = tf.reduce_sum(digits_loss)

    # Calculate the loss from presence indicator being wrong.
    presence_loss = tf.nn.sigmoid_cross_entropy_with_logits(
        logits=y[:, :1], labels=y_[:, :1])
    presence_loss = 7 * tf.reduce_sum(presence_loss)

    return digits_loss, presence_loss, digits_loss + presence_loss
```



CODE

```
def convolutional_layers():
    """
    Get the convolutional layers of the model.
    """

    x = tf.placeholder(tf.float32, [None, None, None])

    # First layer
    W_conv1 = weight_variable([5, 5, 1, 48])
    b_conv1 = bias_variable([48])
    x_expanded = tf.expand_dims(x, 3)
    h_conv1 = tf.nn.relu(conv2d(x_expanded, W_conv1) + b_conv1)
    h_pool1 = max_pool(h_conv1, ksize=(2, 2), stride=(2, 2))

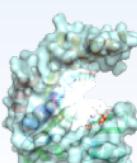
    # Second layer
    W_conv2 = weight_variable([5, 5, 48, 64])
    b_conv2 = bias_variable([64])

    h_conv2 = tf.nn.relu(conv2d(h_pool1, W_conv2) + b_conv2)
    h_pool2 = max_pool(h_conv2, ksize=(2, 1), stride=(2, 1))

    # Third layer
    W_conv3 = weight_variable([5, 5, 64, 128])
    b_conv3 = bias_variable([128])

    h_conv3 = tf.nn.relu(conv2d(h_pool2, W_conv3) + b_conv3)
    h_pool3 = max_pool(h_conv3, ksize=(2, 2), stride=(2, 2))

    return x, h_pool3, [W_conv1, b_conv1,
                        W_conv2, b_conv2,
                        W_conv3, b_conv3]
```



CODE

```
def get_training_model():
    """
    The training model acts on a batch of 128x64 windows, and outputs a
    7 * len(common.CHARS) vector, `v`. `v[0]` is the probability that the
    character is fully within the image and is at the correct scale.

    `v[1 + i * len(common.CHARS) + c]` is the probability that the
    character is `c`.

    ...
    x, conv_layer, conv_vars = convolutional_layers()

    # Densely connected layer
    W_fc1 = weight_variable([32 * 8 * 128, 2048])
    b_fc1 = bias_variable([2048])

    conv_layer_flat = tf.reshape(conv_layer, [-1, 32 * 8 * 128])
    h_fc1 = tf.nn.relu(tf.matmul(conv_layer_flat, W_fc1) + b_fc1)

    # Output layer
    W_fc2 = weight_variable([2048, 1 + 7 * len(common.CHARS)])
    b_fc2 = bias_variable([1 + 7 * len(common.CHARS)])

    y = tf.matmul(h_fc1, W_fc2) + b_fc2

    return (x, y, conv_vars + [W_fc1, b_fc1, W_fc2, b_fc2])
```

```
def get_detect_model():
    """
    The same as the training model, except it acts on an arbitrarily sized
    input, and slides the 128x64 window across the image in 8x8 strides.

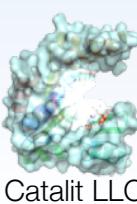
    The output is of the form `v`, where `v[i, j]` is equivalent to the
    output of the training model, for the window at coordinates `(8 * i, 4 * j)`.

    ...
    x, conv_layer, conv_vars = convolutional_layers()

    # Fourth layer
    W_fc1 = weight_variable([8 * 32 * 128, 2048])
    W_conv1 = tf.reshape(W_fc1, [8, 32, 128, 2048])
    b_fc1 = bias_variable([2048])
    h_conv1 = tf.nn.relu(conv2d(conv_layer, W_conv1,
                               stride=(1, 1), padding="VALID") + b_fc1)

    # Fifth layer
    W_fc2 = weight_variable([2048, 1 + 7 * len(common.CHARS)])
    W_conv2 = tf.reshape(W_fc2, [1, 1, 2048, 1 + 7 * len(common.CHARS)])
    b_fc2 = bias_variable([1 + 7 * len(common.CHARS)])
    h_conv2 = conv2d(h_conv1, W_conv2) + b_fc2

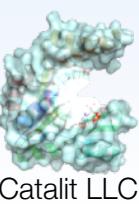
    return (x, h_conv2, conv_vars + [W_fc1, b_fc1, W_fc2, b_fc2])
```



DEEP ANPR SUMMARY

- WHY IS IT IMPORTANT?
 - synthetic data augmentation / generation
 - opencv for image preprocessing and generation
 - separate parallel processes in CPU and GPU
 - convolution on arbitrary size image
 - custom loss
 - different way to save model (older code, no checkpoints yet)

LAB



Catalit LLC